

PHAS0102 - Assignment 4

Student ID: 24022849

Part 1: Time Dependant Problems: implementing finite differences schemes (explicit and implicit) using CPU methods

Introduction / context:

In the context of this problem the von Neumann stability factor (dt/dx^2) was used for our heat equation. The Courant number ($c*dt/dx$) seen in the 1D wave equation in class and the von Neumann stability factor both govern numerical stability but reflect different physical processes. The Courant number ensures numerical waves don't outpace physical wave speed ($CFL \leq 1$), while our heat equation's condition ($dt/dx^2 \leq 0.25$ in 2D) manages diffusive spreading. This fundamental difference appears in their mathematical form: first spatial derivatives for waves versus second spatial derivatives for heat, resulting in distinct stability constraints despite both relating temporal and spatial discretization.

We derive the von Neumann stability factor for this problem by:

1. Using central differences for space (Δu): $\Delta u \approx (u[i+1,j] + u[i-1,j] + u[i,j+1] + u[i,j-1] - 4u[i,j])/dx^2$
2. Taking the Forward Euler for time: $(u^{(n+1)} - u^n)/dt = \Delta u^n$
3. Substituting the discrete Laplacian into the time discretization: $u^{(n+1)} = u^n + (dt/dx^2)(u[i+1,j]^n + u[i-1,j]^n + u[i,j+1]^n + u[i,j-1]^n - 4u[i,j]^n)$
4. For von Neumann stability analysis, let: $u[i,j]^n = \xi^n \exp(i\alpha i + i\beta j)$ where ξ is the amplification factor
5. Substituting: $\xi = 1 + (dt/dx^2)(\exp(i\alpha) + \exp(-i\alpha) + \exp(i\beta) + \exp(-i\beta) - 4) = 1 + (dt/dx^2)(2\cos(\alpha) + 2\cos(\beta) - 4) = 1 - 4(dt/dx^2)(1 - (\cos(\alpha) + \cos(\beta))/2)$
6. For stability need $|\xi| \leq 1$, which leads to: $dt/dx^2 \leq 0.25$

Thus we arrive at the stability condition $dt/dx^2 \leq 0.25$ for the 2D heat equation.

In the following sections a) and b) we build the Explicit (Forward) and Implicit (Backward) schemes respectively and in c) we will build all the functions used to analyse our schemes.

a) Explicit scheme (Forward Euler)

```
import numpy as np
from numba import njit

@njit
def forward_euler_step(u, dx, dt):
```

```

"""
Single step of Forward Euler method
For 2D heat equation:  $du/dt = d^2u/dx^2 + d^2u/dy^2$ 
Forward Euler approximation:  $(u^{(n+1)} - u^{(n)})/dt = (\text{central differences for Laplacian at time } n)$ 

Von Neumann stability analysis requires:  $dt/(dx^2) \leq 0.25$  in 2D
"""
u_new = u.copy()
u_new[1:-1, 1:-1] = u[1:-1, 1:-1] + (dt/dx**2) * (
    u[2:, 1:-1] +      # up
    u[:-2, 1:-1] +     # down
    u[1:-1, 2:] +      # right
    u[1:-1, :-2] -     # left
    4 * u[1:-1, 1:-1] # center
)
return u_new

@njit
def get_center_temp(u, n_points):
    """Get temperature at center point"""
    return u[n_points//2, n_points//2]

def solve_forward_euler(n_points, dt, L):
    """
    Solve 2D heat equation using Forward Euler method until center
    reaches temperature 1.0

    Parameters:
    -----
    n_points : int
        Number of grid points in each dimension
    dt : float
        Time step size (should satisfy  $dt/(dx^2) \leq 0.25$  for stability)
    L : float
        Domain length

    Returns:
    -----
    time : float
        Time taken to reach center temperature of 1.0
    u_initial : ndarray
        Initial temperature field
    u : ndarray
        Final temperature field
    max_values : ndarray
        Maximum temperature at each time step
    """
    dx = L / (n_points - 1)
    n = n_points * n_points

```

```

    stability_factor = dt / (dx * dx)
    print(f"\nForward Euler Solver")
    print(f"Grid size: {n_points}x{n_points}")
    print(f"dx: {dx:.8f}")
    print(f"dt: {dt:.8f}")
    print(f"Forward Euler stability number (dt/dx2):
{stability_factor:.6f}")

    # Initialize solution
    u = np.zeros((n_points, n_points), dtype=np.float32)
    u[:, -1] = 5.0 # Right boundary

    u_initial = u.copy()
    max_values = [np.max(u)]
    physical_time = 0.0

    while get_center_temp(u, n_points) < 1.0:
        u = forward_euler_step(u, dx, dt)
        # Apply boundary conditions
        u[:, -1] = 5.0
        u[0, :] = u[-1, :] = u[:, 0] = 0.0

        physical_time += dt
        max_values.append(np.max(u))

        # Monitor center temperature
        center_temp = get_center_temp(u, n_points)

        # Check for instability
        if np.max(np.abs(u)) > 100:
            print(f"Warning: Forward Euler became unstable at
t={physical_time:.6f}")
            break

    return physical_time, u_initial, u, np.array(max_values),
center_temp

```

b) Implicit scheme (Backward Euler)

```

import numpy as np
from scipy import sparse
from scipy.sparse.linalg import spsolve
from numba import njit
from time import perf_counter

@njit
def assemble_laplacian_numba(n_points, dx):
    """
    Fast matrix assembly using Numba
    Returns arrays for constructing sparse matrix
    """

```

```

"""
n = n_points * n_points
# Pre-calculate number of non-zeros
n_internal = (n_points - 2) * (n_points - 2) # internal points
with 5 entries
n_boundary = 4 * (n_points - 1) # boundary points with 1 entry
nnz = 5 * n_internal + n_boundary

# Pre-allocate arrays
rows = np.zeros(nnz, dtype=np.int32)
cols = np.zeros(nnz, dtype=np.int32)
data = np.zeros(nnz, dtype=np.float64)

idx = 0
for i in range(n_points):
    for j in range(n_points):
        row = i * n_points + j

        if i == 0 or i == n_points-1 or j == 0 or j == n_points-1:
            # Boundary points
            rows[idx] = row
            cols[idx] = row
            data[idx] = 1.0
            idx += 1
        else:
            # Interior points - 5-point stencil
            # Center
            rows[idx] = row
            cols[idx] = row
            data[idx] = -4.0
            idx += 1

            # Up
            rows[idx] = row
            cols[idx] = row - n_points
            data[idx] = 1.0
            idx += 1

            # Down
            rows[idx] = row
            cols[idx] = row + n_points
            data[idx] = 1.0
            idx += 1

            # Left
            rows[idx] = row
            cols[idx] = row - 1
            data[idx] = 1.0
            idx += 1

```

```

        # Right
        rows[idx] = row
        cols[idx] = row + 1
        data[idx] = 1.0
        idx += 1

    return rows, cols, data / (dx * dx)

def create_2d_laplacian(n_points, dx):
    """
    Create sparse Laplacian matrix using Numba-accelerated assembly
    """
    n = n_points * n_points
    rows, cols, data = assemble_laplacian_numba(n_points, dx)
    return sparse.csr_matrix((data, (rows, cols)), shape=(n, n))

@njit
def get_center_temp(u, n_points):
    """Get temperature at center point"""
    return u[n_points//2, n_points//2]

def solve_backward_euler(n_points, dt, L):
    """
    Solve 2D heat equation using Backward Euler method until center
    reaches temperature 1.0

    Parameters:
    -----
    n_points : int
        Number of grid points in each dimension
    dt : float
        Time step size (unconditionally stable, but larger dt affects
accuracy)
    L : float
        Domain length

    Returns:
    -----
    time : float
        Time taken to reach center temperature of 1.0
    u_initial : ndarray
        Initial temperature field
    u : ndarray
        Final temperature field
    max_values : ndarray
        Maximum temperature at each time step
    center_temp : float
        Final temperature at center point
    """
    dx = L / (n_points - 1)

```

```

n = n_points * n_points
stability_factor = dt / (dx * dx)
print(f"\nBackward Euler Solver")
print(f"Grid size: {n_points}x{n_points}")
print(f"dx: {dx:.8f}")
print(f"dt: {dt:.8f}")
print(f"Backward Euler stability number (dt/dx²):
{stability_factor:.6f}")

# Create system matrices
start_time = perf_counter()

# Create the Laplacian matrix
A = create_2d_laplacian(n_points, dx)
I = sparse.eye(n, format='csr')

# Create implicit system matrix: (I - dt*A)
A_implicit = (I - dt * A).tocsr()

# Modify matrix for boundary conditions
for i in range(n_points):
    for j in range(n_points):
        if i == 0 or i == n_points-1 or j == 0 or j == n_points-1:
            row = i * n_points + j

A_implicit.data[A_implicit.indptr[row]:A_implicit.indptr[row+1]] = 0.0
A_implicit[row, row] = 1.0

A_implicit.eliminate_zeros()
print(f"Matrix setup time: {perf_counter() - start_time:.2f}
seconds")

# Initialize solution
u = np.zeros((n_points, n_points), dtype=np.float64)
u[:, -1] = 5.0 # Right boundary

u_initial = u.copy()
max_values = [np.max(u)]

# Time stepping loop
physical_time = 0.0
solve_start = perf_counter()
steps = 0

while get_center_temp(u, n_points) < 1.0:
    # Create right-hand side
    b = u.reshape(-1).copy()

    # Set boundary values in right-hand side
    for i in range(n_points):

```

```

        for j in range(n_points):
            idx = i * n_points + j
            if i == 0 or i == n_points-1 or j == 0:
                b[idx] = 0.0
            elif j == n_points-1:
                b[idx] = 5.0

    # Solve system
    x = spsolve(A_implicit, b)
    u = x.reshape((n_points, n_points))

    # Update time
    physical_time += dt
    max_values.append(np.max(u))
    steps += 1

    # Monitor center temperature
    center_temp = get_center_temp(u, n_points)

    if steps % 100 == 0:
        elapsed = perf_counter() - solve_start
        print(f"Step {steps}, Time {physical_time:.8f}, "
              f"Center temp: {center_temp:.8f}")

    return physical_time, u_initial, u, np.array(max_values),
    center_temp

```

c) Create set of functions used for:

1. Checking and plotting the convergence of both schemes to $t^* = 0.424011387033$ as the number of discretisation points increases.
2. Numerically investigate the stability of both schemes: we will investigate a set of stability factors (which determine the size of the time step) and their impact on the simulation, plot the maximum temperature reached for each factor and add some visuals of the plate in its initial state and final state (i.e. when the temperature at the center reaches 1) for each factor analysed.

```

import numpy as np
import matplotlib.pyplot as plt
from time import time as current_time

# Physical parameters and configuration common for all functions
L = 2 # Domain length [-1, 1] -> length = 2
t_exact = 0.424011387033 # Known exact time for center temperature = 1

# Convergence study function
def run_convergence_study(grid_sizes, stability_factor_forward,
                           stability_factor_backward):
    """Run and plot convergence study for both methods using

```

```

appropriate timesteps"""
    results = {
        "forward": {'times': [], 'errors': [], 'compute_times': [],
'final_temperature': []},
        "backward": {'times': [], 'errors': [], 'compute_times': [],
'final_temperature': []}
    }

    for grid_size in grid_sizes:
        dx = L / (grid_size - 1)

        # Forward Euler: use conservative timestep
        dt_forward = stability_factor_forward * dx * dx
        start_time = current_time()
        t_final, _, _, _, center_temp = solve_forward_euler(grid_size,
dt_forward, L)
        compute_time = current_time() - start_time
        error = abs(t_exact - t_final)
        results["forward"][['times']].append(t_final)
        results["forward"][['errors']].append(error)
        results["forward"][['compute_times']].append(compute_time)
        results["forward"][['final_temperature']].append(center_temp)
        print(f"Forward Euler - Grid: {grid_size}, Center Temperature
Result: {center_temp:.6f}, Physical Time Result: {t_final:.6f}, Error:
{error:.6f}, Compute Time: {compute_time:.3f}s")

        # Backward Euler: can use larger timestep
        dt_backward = stability_factor_backward * dx * dx
        start_time = current_time()
        t_final, _, _, _, center_temp =
solve_backward_euler(grid_size, dt_backward, L)
        compute_time = current_time() - start_time
        error = abs(t_exact - t_final)
        results["backward"][['times']].append(t_final)
        results["backward"][['errors']].append(error)
        results["backward"][['compute_times']].append(compute_time)
        results["backward"][['final_temperature']].append(center_temp)
        print(f"Backward Euler - Grid: {grid_size}, Center Temperature
Result: {center_temp:.6f}, Physical Time Result: {t_final:.6f}, Error:
{error:.6f}, Compute Time: {compute_time:.3f}s")

    # Plot results
    fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(18, 5))

    # Convergence plot
    for method in results:
        ax1.plot(grid_sizes, results[method][['times']], marker='o',
label=f"{method.capitalize()} Euler")
        ax1.axhline(y=t_exact, color='k', linestyle='--', label='t* =
0.42401138703')

```



```

ax1.set_xlabel("Grid size")
ax1.set_ylabel("Time to reach center temperature = 1")
ax1.set_title("Convergence Comparison")
ax1.legend()
ax1.grid(True)

# Error plot
for method in results:
    ax2.plot(grid_sizes, results[method]['errors'], marker='o',
             label=f"{method.capitalize()} Euler")
ax2.set_xlabel("Grid size")
ax2.set_ylabel("Error")
ax2.set_title("Error Analysis")
ax2.set_yscale('log')
ax2.legend()
ax2.grid(True)

# Computation time plot
for method in results:
    ax3.plot(grid_sizes, results[method]['compute_times'],
marker='o',
             label=f"{method.capitalize()} Euler")
ax3.set_xlabel("Grid size")
ax3.set_ylabel("Computation time (s)")
ax3.set_title("Performance Comparison")
ax3.set_yscale('log')
ax3.legend()
ax3.grid(True)

plt.tight_layout()
plt.show()

# Stability Analysis function
def run_stability_analysis(grid_size, stability_factors_forward,
stability_factors_backward):
    """Run and plot stability analysis showing different stability
    properties"""
    grid_size = grid_size
    dx = L / (grid_size - 1)

    # Create figure with two subplots side by side
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

    # Forward Euler plot
    for st_factor in stability_factors_forward:
        dt = st_factor * dx * dx
        tfinal, _, _, maxvalues, _ = solve_forward_euler(grid_size,
dt, L)
        timesteps = np.arange(len(maxvalues)) * dt

```

```

        ax1.plot(timesteps, maxvalues,
                  label=f' $dt/dx^2 = \{st\_factor:.3f\}$ ')

    ax1.set_xlabel('Time')
    ax1.set_ylabel('Maximum Temperature')
    ax1.set_title(f'Forward Euler\nGrid Size {grid_size}x{grid_size}')
    ax1.legend()
    ax1.grid(True)

    # Backward Euler plot
    for st_factor in stability_factors_backward:
        dt = st_factor * dx * dx
        tfinal, _, _, maxvalues, _ = solve_backward_euler(grid_size,
dt, L)
        timesteps = np.arange(len(maxvalues)) * dt

        ax2.plot(timesteps, maxvalues,
                  label=f' $dt/dx^2 = \{st\_factor:.3f\}$ ')

    ax2.set_xlabel('Time')
    ax2.set_ylabel('Maximum Temperature')
    ax2.set_title(f'Backward Euler\nGrid Size
{grid_size}x{grid_size}')
    ax2.legend()
    ax2.grid(True)

    plt.tight_layout()
    plt.show()

# Solution fields function
def plot_solution_fields(grid_size, stability_factor_forward,
stability_factor_backward):
    """Plot initial and final temperature fields for both methods"""
    grid_size = grid_size
    dx = L / (grid_size - 1)

    # Forward Euler
    dt_forward = stability_factor_forward * dx * dx
    _, initial_f, final_f, _, _ = solve_forward_euler(grid_size,
dt_forward, L)

    # Backward Euler
    dt_backward = stability_factor_backward * dx * dx
    _, initial_b, final_b, _, _ = solve_backward_euler(grid_size,
dt_backward, L)

    # Plot all results in a 2x2 grid
    fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(12,
10))

```

```

# Forward Euler results
im1 = ax1.imshow(initial_f, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(im1, ax=ax1, label='Temperature')
ax1.set_title(f"Initial Condition: N = {grid_size}\nForward Euler
(dt/dx2 = {stability_factor_forward})")
ax1.set_xlabel('x')
ax1.set_ylabel('y')

im2 = ax2.imshow(final_f, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(im2, ax=ax2, label='Temperature')
ax2.set_title(f"Final Condition: N = {grid_size}\nForward Euler
(dt/dx2 = {stability_factor_forward})")
ax2.set_xlabel('x')
ax2.set_ylabel('y')

# Backward Euler results
im3 = ax3.imshow(initial_b, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(im3, ax=ax3, label='Temperature')
ax3.set_title(f"Initial Condition: N = {grid_size}\nBackward Euler
(dt/dx2 = {stability_factor_backward})")
ax3.set_xlabel('x')
ax3.set_ylabel('y')

im4 = ax4.imshow(final_b, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(im4, ax=ax4, label='Temperature')
ax4.set_title(f"Final Condition: N = {grid_size}\nBackward Euler
(dt/dx2 = {stability_factor_backward})")
ax4.set_xlabel('x')
ax4.set_ylabel('y')

plt.tight_layout()
plt.show()

```

c1) Convergence analysis

We run both schemes for different grid sizes. The choice of the stability factor has several trade offs, some of which are described below.

Forward Euler:

- Must keep $dt/(dx^2) \leq 0.25$ for stability in 2D in a heat equation ($u_t = \Delta u$)
- Smaller values = more accurate but more time steps needed
- Larger values (approaching 0.25) = faster computation but risk instability
- If $dt/(dx^2) > 0.25$: solution will become unstable and oscillate

Using the Explicit method means each step is computationally cheap, but the stability restriction means many small steps required.

Backward Euler:

- Unconditionally stable for any $dt/(dx^2)$
- Each step more expensive (must solve linear system)
- The smaller the stability factor $dt/(dx^2)$ the more accurate the scheme will be (closer to the "true" solution), but more time steps will be required making it computationally more expensive

In this exercise we are running the Forward Euler scheme at the stability limit (i.e. $dt/(dx^2) = 0.25$) and the Backward Euler scheme at 16. These are relatively large factors but will allow to run several grid sizes in acceptable times and still achieve a decent level of precision.

```
from google.colab import output
output.no_vertical_scroll()

# Grid configurations
grid_sizes = [100, 150, 250, 300] # Main grid sizes for convergence study

# Set stability factor
stability_factor_forward = 0.25 # Stability factor for Forward Euler
stability_factor_backward = 16 # Stability factor for Backward Euler

run_convergence_study(grid_sizes, stability_factor_forward,
stability_factor_backward)

<IPython.core.display.Javascript object>
```

Forward Euler Solver

Grid size: 100x100

dx: 0.02020202

dt: 0.00010203

Forward Euler stability number (dt/dx^2): 0.250000

Forward Euler - Grid: 100, Center Temperature Result: 1.000021,

Physical Time Result: 0.407713, Error: 0.016298, Compute Time: 2.395s

Backward Euler Solver

Grid size: 100x100

dx: 0.02020202

dt: 0.00652995

Backward Euler stability number (dt/dx^2): 16.000000

Matrix setup time: 0.86 seconds

Backward Euler - Grid: 100, Center Temperature Result: 1.004681,

Physical Time Result: 0.417917, Error: 0.006095, Compute Time: 3.701s

Forward Euler Solver

Grid size: 150x150

dx: 0.01342282
dt: 0.00004504
Forward Euler stability number (dt/dx^2): 0.250000
Forward Euler - Grid: 150, Center Temperature Result: 1.000048,
Physical Time Result: 0.413090, Error: 0.010922, Compute Time: 3.810s

Backward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00288275
Backward Euler stability number (dt/dx^2): 16.000000
Matrix setup time: 0.04 seconds
Step 100, Time 0.28827530, Center temp: 0.77095928
Backward Euler - Grid: 150, Center Temperature Result: 1.002541,
Physical Time Result: 0.417999, Error: 0.006012, Compute Time: 18.277s

Forward Euler Solver
Grid size: 250x250
dx: 0.00803213
dt: 0.00001613
Forward Euler stability number (dt/dx^2): 0.250000
Forward Euler - Grid: 250, Center Temperature Result: 1.000021,
Physical Time Result: 0.417413, Error: 0.006599, Compute Time: 29.636s

Backward Euler Solver
Grid size: 250x250
dx: 0.00803213
dt: 0.00103224
Backward Euler stability number (dt/dx^2): 16.000000
Matrix setup time: 0.36 seconds
Step 100, Time 0.10322414, Center temp: 0.13936125
Step 200, Time 0.20644828, Center temp: 0.53209932
Step 300, Time 0.30967242, Center temp: 0.81735081
Step 400, Time 0.41289657, Center temp: 0.99281241
Backward Euler - Grid: 250, Center Temperature Result: 1.000785,
Physical Time Result: 0.419090, Error: 0.004921, Compute Time:
251.327s


Forward Euler Solver
Grid size: 300x300
dx: 0.00668896
dt: 0.00001119
Forward Euler stability number (dt/dx^2): 0.250000
Forward Euler - Grid: 300, Center Temperature Result: 1.000007,
Physical Time Result: 0.418496, Error: 0.005515, Compute Time: 50.354s

Backward Euler Solver
Grid size: 300x300
dx: 0.00668896
dt: 0.00071588

```

Backward Euler stability number (dt/dx2): 16.000000
Matrix setup time: 0.10 seconds
Step 100, Time 0.07158757, Center temp: 0.04338260
Step 200, Time 0.14317513, Center temp: 0.29337280
Step 300, Time 0.21476270, Center temp: 0.55933020
Step 400, Time 0.28635026, Center temp: 0.76327129
Step 500, Time 0.35793783, Center temp: 0.90948389
Backward Euler - Grid: 300, Center Temperature Result: 1.000336,
Physical Time Result: 0.419503, Error: 0.004508, Compute Time:
494.377s

```



c2) Numerical investigation of the stability of both schemes

In this analysis we will fix the grid size at 150x150 (choosing this size so that solving and plotting does not take too long) and run both schemes with different stability factors.

For the Forward Euler we will run some below the 0.25 limit and some above to see if the solution does keep stability below the limit and loses it above the limit.

We do not expect the Backward Euler method to lose stability at any factor, so we will start with a factor slightly below the one used in the convergence analysis and scale it up a fair bit to see if stability is maintained.

The script checks the maximum temperature values across the grid and how they change when running with different stability factors.

```

from google.colab import output
output.no_vertical_scroll()

# Fixed grid size for stability factor comparison
grid_size = grid_sizes[1]

stability_factors_forward = [0.1, 0.2, 0.25, 0.5, 0.75, 1] # first
# three within stability limit, last three outside
stability_factors_backward = [8.0, 12.0, 16.0, 32.0, 64.0, 128.0] #
# Can use much larger steps
run_stability_analysis(grid_size, stability_factors_forward,
stability_factors_backward)

```

<IPython.core.display.Javascript object>

Forward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00001802
Forward Euler stability number (dt/dx^2): 0.100000

Forward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00003603
Forward Euler stability number (dt/dx^2): 0.200000

Forward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00004504
Forward Euler stability number (dt/dx^2): 0.250000

Forward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00009009
Forward Euler stability number (dt/dx^2): 0.500000
Warning: Forward Euler became unstable at $t=0.000901$

Forward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00013513
Forward Euler stability number (dt/dx^2): 0.750000
Warning: Forward Euler became unstable at $t=0.000811$

Forward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00018017
Forward Euler stability number (dt/dx^2): 1.000000
Warning: Forward Euler became unstable at $t=0.000901$

Backward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00144138
Backward Euler stability number (dt/dx^2): 8.000000
Matrix setup time: 0.06 seconds
Step 100, Time 0.14413765, Center temp: 0.30140223
Step 200, Time 0.28827530, Center temp: 0.77328956

Backward Euler Solver


Grid size: 150x150
dx: 0.01342282
dt: 0.00216206
Backward Euler stability number (dt/dx^2): 12.000000
Matrix setup time: 0.04 seconds
Step 100, Time 0.21620648, Center temp: 0.56794706

Backward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00288275
Backward Euler stability number (dt/dx^2): 16.000000
Matrix setup time: 0.04 seconds
Step 100, Time 0.28827530, Center temp: 0.77095928

Backward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.00576551
Backward Euler stability number (dt/dx^2): 32.000000
Matrix setup time: 0.06 seconds

Backward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.01153101
Backward Euler stability number (dt/dx^2): 64.000000
Matrix setup time: 0.04 seconds

Backward Euler Solver
Grid size: 150x150
dx: 0.01342282
dt: 0.02306202
Backward Euler stability number (dt/dx^2): 128.000000
Matrix setup time: 0.04 seconds



The above plots show that the maximum temperature across the grid during our simulations never exceeds 5 for the Backward Euler scheme. However, for the Forward Euler scheme we can see that some simulations "exploded". This is not very clear in a line graph since several lines are likely overlapping.

We will look at them individually by plotting the 2D heat map for each factor.

We've plotted our set of stability factors in pairs, one per method.

```
from google.colab import output
output.no_vertical_scroll()
import warnings
warnings.filterwarnings('ignore',
category=sparse.SparseEfficiencyWarning) # ignoring warnings in order
to keep result space readable

# Fixed grid size for stability factor comparison
grid_size = grid_sizes[1]

# Timestep factors for each method based on their stability
properties:
# Forward Euler: must have  $dt/(dx^2) \leq 1/4$  for stability in 2D

stability_factors_forward = [0.1, 0.2, 0.25, 0.5, 0.75, 1] # first
three within stability limit, last three outside
stability_factors_backward = [8.0, 12.0, 16.0, 32.0, 64.0, 128.0] #
Can use much larger steps

for sff, sfb in zip(stability_factors_forward,
stability_factors_backward):
    plot_solution_fields(grid_size, stability_factor_forward=sff,
stability_factor_backward=sfb)

<IPython.core.display.Javascript object>
```

Forward Euler Solver

Grid size: 150x150

dx: 0.01342282

dt: 0.00001802

Forward Euler stability number (dt/dx^2): 0.100000

Backward Euler Solver

Grid size: 150x150

dx: 0.01342282



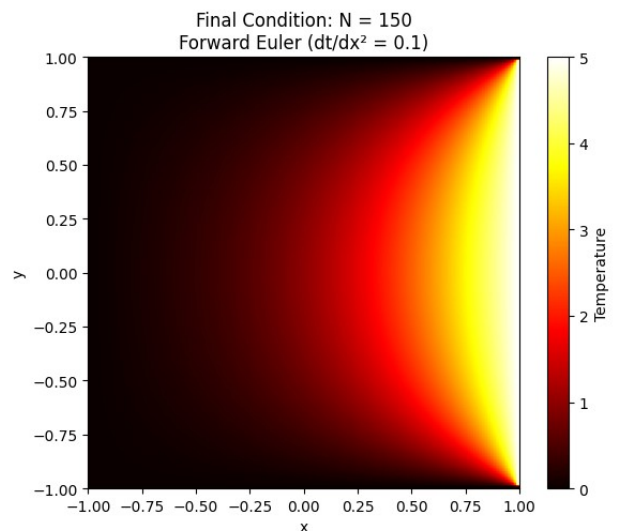

dt: 0.00144138

Backward Euler stability number (dt/dx^2): 8.000000

Matrix setup time: 0.04 seconds

Step 100, Time 0.14413765, Center temp: 0.30140223

Step 200, Time 0.28827530, Center temp: 0.77328956



Forward Euler Solver

Grid size: 150x150

dx: 0.01342282

dt: 0.00003603

Forward Euler stability number (dt/dx^2): 0.200000

Backward Euler Solver

Grid size: 150x150


dx: 0.01342282

dt: 0.00216206

Backward Euler stability number (dt/dx^2): 12.000000

Matrix setup time: 0.04 seconds

Step 100, Time 0.21620648, Center temp: 0.56794706


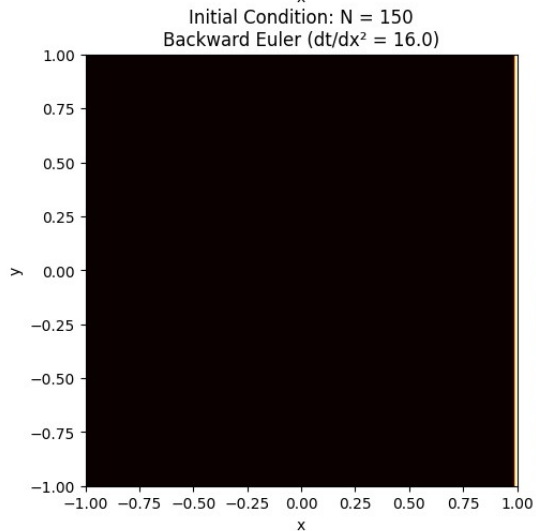

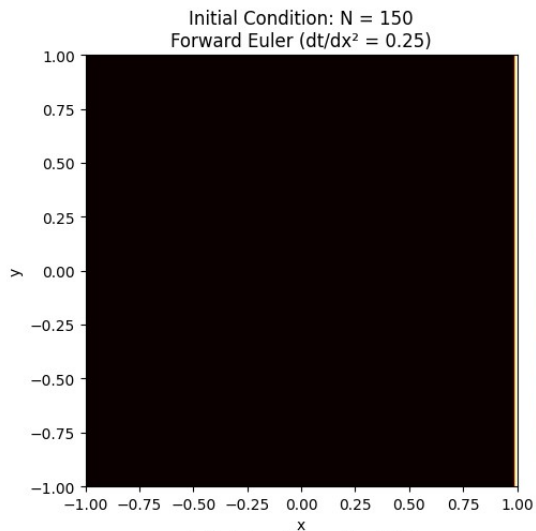


Forward Euler Solver

Grid size: 150x150
dx: 0.01342282
dt: 0.00004504
Forward Euler stability number (dt/dx^2): 0.250000

Backward Euler Solver

Grid size: 150x150
dx: 0.01342282
dt: 0.00288275
Backward Euler stability number (dt/dx^2): 16.000000
Matrix setup time: 0.07 seconds
Step 100, Time 0.28827530, Center temp: 0.77095928



Forward Euler Solver
Grid size: 150x150
dx: 0.01342282

dt: 0.00009009
Forward Euler stability number (dt/dx^2): 0.500000
Warning: Forward Euler became unstable at $t=0.000901$

Backward Euler Solver


Grid size: 150x150

dx: 0.01342282

dt: 0.00576551

Backward Euler stability number (dt/dx^2): 32.000000

Matrix setup time: 0.06 seconds



Forward Euler Solver

Grid size: 150x150

dx: 0.01342282

dt: 0.00013513

Forward Euler stability number (dt/dx^2): 0.750000

Warning: Forward Euler became unstable at $t=0.000811$

Backward Euler Solver


Grid size: 150x150

dx : 0.01342282

dt : 0.01153101

Backward Euler stability number (dt/dx^2): 64.000000

Matrix setup time: 0.06 seconds



Forward Euler Solver

Grid size: 150x150

dx : 0.01342282

dt : 0.00018017

Forward Euler stability number (dt/dx^2): 1.000000

Warning: Forward Euler became unstable at $t=0.000901$

Backward Euler Solver


Grid size: 150x150

dx: 0.01342282

dt: 0.02306202

Backward Euler stability number (dt/dx^2): 128.000000

Matrix setup time: 0.07 seconds



One thing that is immediately visible for all simulations is that the right hand side is heated up to 5 degrees in all initial states (small white vertical stripe on the righthand side, which will be wider if the resolution of the grid is lower and vice versa) as indicated in the problem's requirements.

By looking at the above plots we can confirm that the Backward Euler scheme maintains stability regardless of what factor (and, consequently what the size of the time step is) while the Forward Euler scheme becomes unstable as soon as we go above our stability factor limit of 0.25.

Part 2: Time Dependant Problems: implementing finite differences explicit scheme using GPU methods

a) Setting up the Forward Euler solver adapted for GPU

```
import numpy as np
import matplotlib.pyplot as plt
from numba import cuda
import math

# GPU kernel for explicit time-stepping
@cuda.jit
def forward_euler_gpu(u, u_new, N, dt, dx2, center_value):
    idx = cuda.grid(1)
    if idx < N * N:
        i = idx // N # Row index
        j = idx % N # Column index

        # Skip boundary points
        if i == 0 or i == N - 1 or j == 0 or j == N - 1:
            u_new[idx] = u[idx]
        else:
            # Compute Laplacian using neighbours
            laplacian = (u[idx - 1] + u[idx + 1] + u[idx - N] + u[idx + N] - 4 * u[idx]) / dx2
            u_new[idx] = u[idx] + dt * laplacian

        # Update the center value (shared memory or global)
        if idx == (N // 2) * N + (N // 2):
            center_value[0] = u_new[idx]

# GPU kernel for initialization
@cuda.jit
def initialize_boundary_conditions(u, N):
    idx = cuda.grid(1)
    if idx < N * N:
        i = idx // N
        j = idx % N
        if j == N - 1:
            u[idx] = 5
        elif i == 0 or i == N - 1 or j == 0:
            u[idx] = 0

# GPU accelerated heat equation solver
def solve_heat_equation_gpu(N, dt, method="Forward Euler"):
    dx = L / (N - 1)
    dx2 = dx ** 2

    # Initialize arrays
    u = np.zeros(N * N, dtype=np.float32)
```



```

u_new = np.zeros_like(u)

# Apply boundary conditions
threads_per_block = 256
blocks_per_grid = math.ceil(u.size / threads_per_block)
initialize_boundary_conditions[blocks_per_grid, threads_per_block](
    u, N)

# Synchronize and copy back to the host to verify correctness
cuda.synchronize()
initial_condition = u.reshape((N, N)).copy()

# Save initial condition before iteration
initial_condition = u.reshape((N, N)).copy()

# Copy to device
d_u = cuda.to_device(u)
d_u_new = cuda.to_device(u_new)
d_center_value = cuda.to_device(np.array([0], dtype=np.float32))
# Device variable for center value

time = 0

while True:
    forward_euler_gpu[blocks_per_grid, threads_per_block](d_u,
d_u_new, N, dt, dx2, d_center_value)
    cuda.synchronize() # Ensure kernel completion
    d_u, d_u_new = d_u_new, d_u # Swap the device arrays

    # Check center value without copying entire array
    center_value = d_center_value.copy_to_host()[0]
    if center_value >= 1:
        break

    time += dt

# Final condition after iteration
final_condition = d_u.copy_to_host().reshape((N, N)).copy()

return time, initial_condition, final_condition, center_value

```

b) Running the solver. Since we are running on GPU we will increase the grid sizes used and reduce the stability factor to 0.1, implying a smaller time step for each grid size.

```

import numpy as np
import matplotlib.pyplot as plt
from numba import cuda, config
import math
from time import time as current_time
from google.colab import output

```

```

output.no_vertical_scroll()

config.CUDA_LOW_OCCUPANCY_WARNINGS = False

# Parameters
L = 2
t_exact = 0.42401138703
stability_factor_forward_gpu = 0.15

# Grid sizes
grid_sizes = [150, 300, 500, 1000, 1500]

# Supporting variables
convergence_results_gpu = []
all_initial_conditions = []
final_conditions = []
errors = []

for grid_size in grid_sizes:
    dx = L / (grid_size - 1)
    dt_forward_gpu = stability_factor_forward_gpu * dx * dx

    print(f"\nForward Euler Solver - GPU")
    print(f"Grid size: {grid_size}x{grid_size}")
    print(f"dx: {dx:.8f}")
    print(f"dt: {dt_forward_gpu:.8f}")
    print(f"Forward Euler stability number (dt/dx²):
{stability_factor_forward_gpu:.6f}")

    start_time = current_time()

    t_gpu, initial_condition, final_condition, center_value =
solve_heat_equation_gpu(grid_size, dt_forward_gpu, method="Forward
Euler")

    convergence_results_gpu.append((grid_size, t_gpu))
    all_initial_conditions.append((grid_size, initial_condition))
    final_conditions.append((grid_size, final_condition))

    error = abs(t_exact - t_gpu)
    errors.append(error)

    compute_time = current_time() - start_time

    print(f"Forward Euler - Grid: {grid_size}, Center Temperature
Result: {center_value:.6f}, Physical Time Result: {t_gpu:.6f}, Error:
{error:.6f}, Compute Time: {compute_time:.3f}s")

# Plot GPU results as a line graph
plt.figure()

```

```

grid_sizes, times = zip(*convergence_results_gpu)
plt.plot(grid_sizes, times, marker='o', label="GPU Convergence
Results")
plt.axhline(y=t_exact, color='k', linestyle='--', label=f't* =
{t_exact}')
plt.xlabel("Grid size")
plt.ylabel("Time to reach center temperature = 1")
plt.title("GPU Convergence Results for Forward Euler")
plt.legend()
plt.grid()

# Plot error vs grid sizes
plt.figure()
plt.plot(grid_sizes, errors, marker='o', label="Error (|t* -
result|)")
plt.xlabel("Grid size")
plt.ylabel("Error")
plt.title("Error Analysis for GPU Convergence Results")
plt.legend()
plt.grid()
plt.tight_layout()
plt.show()

# Plot initial conditions before iteration for the first grid size
initial_grid_size, initial_condition = all_initial_conditions[0]
plt.figure()
plt.imshow(initial_condition, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(label='Temperature')
plt.title(f"Initial Condition (Before Iteration) for Grid Size
{initial_grid_size}")
plt.xlabel('x')
plt.ylabel('y')

# Plot the grid after iteration for the first grid size
final_grid_size, final_condition = final_conditions[0]
plt.figure()
plt.imshow(final_condition, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(label='Temperature')
plt.title(f"Final Condition (After Iteration) for Grid Size
{final_grid_size}")
plt.xlabel('x')
plt.ylabel('y')
plt.tight_layout()
plt.show()

# Plot initial conditions before iteration for the last grid size
initial_grid_size, initial_condition = all_initial_conditions[-1]
plt.figure()

```

```

plt.imshow(initial_condition, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(label='Temperature')
plt.title(f"Initial Condition (Before Iteration) for Grid Size
{initial_grid_size}")
plt.xlabel('x')
plt.ylabel('y')

# Plot the grid after iteration for the last grid size
final_grid_size, final_condition = final_conditions[-1]
plt.figure()
plt.imshow(final_condition, extent=[-1, 1, -1, 1], origin='lower',
cmap='hot')
plt.colorbar(label='Temperature')
plt.title(f"Final Condition (After Iteration) for Grid Size
{final_grid_size}")
plt.xlabel('x')
plt.ylabel('y')
plt.tight_layout()
plt.show()

```

```

print("GPU Convergence Results with Errors:")
for i, (grid_size, t_gpu) in enumerate(convergence_results_gpu):
    print(f"Grid size: {grid_size}, Result: {t_gpu}, Error:
{errors[i]}")

```

<IPython.core.display.Javascript object>

Forward Euler Solver - GPU

Grid size: 150x150

dx: 0.01342282

dt: 0.00002703

Forward Euler stability number (dt/dx^2): 0.150000

/usr/local/lib/python3.10/dist-packages/numba/cuda/cudadrv/
devicearray.py:888: NumbaPerformanceWarning: Host array used in CUDA
kernel will incur copy overhead to/from device.

warn(NumbaPerformanceWarning(msg))

Forward Euler - Grid: 150, Center Temperature Result: 1.000010,

Physical Time Result: 0.413062, Error: 0.010949, Compute Time: 5.645s

Forward Euler Solver - GPU

Grid size: 300x300

dx: 0.00668896

dt: 0.00000671

Forward Euler stability number (dt/dx^2): 0.150000

Forward Euler - Grid: 300, Center Temperature Result: 1.000005,

Physical Time Result: 0.418512, Error: 0.005499, Compute Time: 10.847s

Forward Euler Solver - GPU

Grid size: 500x500

dx: 0.00400802

dt: 0.00000241

Forward Euler stability number (dt/dx^2): 0.150000

Forward Euler - Grid: 500, Center Temperature Result: 1.000001,

Physical Time Result: 0.420774, Error: 0.003237, Compute Time: 38.992s

Forward Euler Solver - GPU

Grid size: 1000x1000

dx: 0.00200200

dt: 0.00000060

Forward Euler stability number (dt/dx^2): 0.150000

Forward Euler - Grid: 1000, Center Temperature Result: 1.000001,

Physical Time Result: 0.422417, Error: 0.001594, Compute Time:
194.378s

Forward Euler Solver - GPU

Grid size: 1500x1500


dx: 0.00133422

dt: 0.00000027


Forward Euler stability number (dt/dx^2): 0.150000


Forward Euler - Grid: 1500, Center Temperature Result: 1.000000,


Physical Time Result: 0.424399, Error: 0.000388, Compute Time:
743.990s

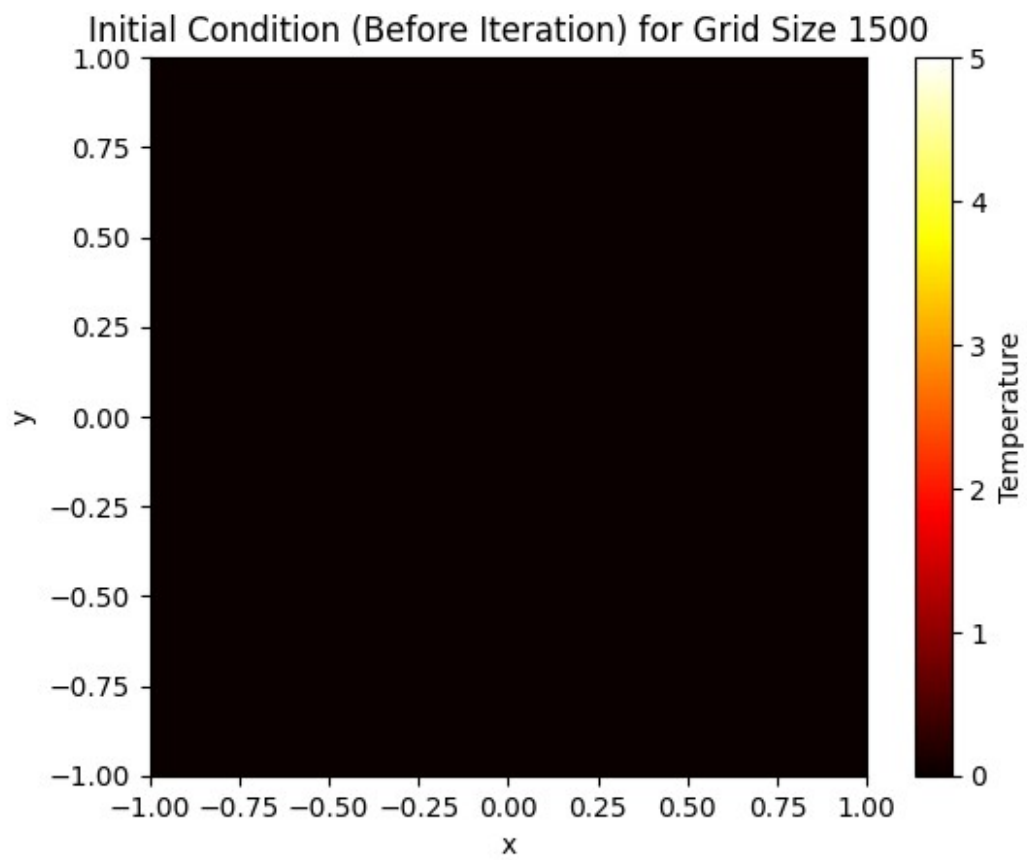



Error Analysis for GPU Convergence Results











GPU Convergence Results with Errors:

Grid size: 150, Result: 0.41306247466322804, Error: 0.010948912366771979

Grid size: 300, Result: 0.4185120971799401, Error: 0.00549928985005993

Grid size: 500, Result: 0.4207742137573994, Error: 0.0032371732726006464

Grid size: 1000, Result: 0.422417011602732, Error: 0.0015943754272680355

Grid size: 1500, Result: 0.4243992770918108, Error: 0.00038789006181078056

Our GPU results clearly benefit from the higher grid resolution and smaller time steps used (implied by using a smaller stability factor).

The error behaves as in the CPU version, getting smaller as the grid size increases.

At $N = 1500$ and $dt/(dx^2) = 0.15$ we even achieve precision of 3 digits (result of 0.4243992770918108) against our target time. One noticeable fact is that our physical time in this simulation actually goes over the target time, unlike simulations at smaller grid sizes. This indicates that at this grid size we would likely benefit from a small stability factor (and hence smaller time step size).

In addition, the heat maps of the plate at $N = 1500$ have so much resolution that the vertical stripe on the righthand side at temperature = 5 is not even visible.